

A Step-by-Step Guide to Pandas Pivot Tables

The Ultimate Guide for Python and Pandas



Let's get started! Source: Nik Piepenbreier

You may be familiar with pivot tables in Excel to generate easy insights into your data. In this post, we'll explore how to create Python pivot tables using the *pivot table* function available in Pandas. The function itself is quite easy to use, but it's not the most intuitive. This post will give you a complete overview of how to use the function!

The function is quite similar to the group by function also available in Pandas, but offers significantly more customization, as we'll see later on in this post. If you

thought this post was about undoing your hard work to pivot your data, check out this [tutorial on un-pivoting your data](#)!

If you want to follow along with a video tutorial, check out my tutorial here:

Python Pivot Tables. Source: Nik Piepenbreier

How to Build a Pivot Table in Python

First off, let's quickly cover off what a pivot table actually is: it's a table of statistics that helps summarize the data of a larger table by "pivoting" that data. In Pandas, we can construct a pivot table using the following syntax, as described in the [official Pandas documentation](#):

The function returns its own dataframe that can be accessed similar to any other dataframe you may come across. Let's take a moment to explore the different parameters available in the function:

Now that we have an understanding of the different parameters available in the function, let's load in our data set and begin exploring our data.

Let's start off by loading our data set! We've built a free downloadable data set that can be found at this link. We'll use Pandas to import the data into a dataframe called **df**. We'll also print out the first five rows using the **.head()** function:

Out:

	Date	Region	Type	Units
0	2020-07-11	East	Children's Clothing	18
1	2020-09-23	North	Children's Clothing	14
2	2020-04-02	South	Women's Clothing	17
3	2020-02-28	East	Children's Clothing	26
4	2020-03-19	West	Women's Clothing	3

Based on the output of the first five rows shown above, we can see that we have five columns to work with:

Now that we have a bit more context around the data, let's explore creating our first pivot table in Pandas.

Creating a Pivot Table in Pandas

To get started with creating a pivot table in Pandas, let's build a very simple pivot table to start things off. We'll begin by aggregating the Sales values by the Region the sale took place in:

This returns the following output:

Sales	Region
East	408.182482
North	438.924051
South	432.956204
West	452.029412

This gave us a summary of the **Sales** field by **Region**.

Recall from our discussion around parameters earlier, that the default parameter for **aggfunc** is **mean**. Because of this, the **Sales** field in the resulting dataframe is the average of **Sales** per **Region**.

If we wanted to change the type of function used, we could use the **aggfunc** parameter. For example, if we wanted to return the **sum** of all **Sales** across a **region**, we could write:

This returns:

Sales	Region
East	167763
North	138700
South	59315
West	61476

Multi-Index Dataframes in Pandas

We can already notice a difference between the dataframe that this function put out, compared to the original dataframe (**df**) we put together. This is because the resulting dataframe from a pivot table function is a MultiIndex dataframes. You can learn more about these dataframes [at this link](#).

As this dataframe is a very simple dataframe, we can simply reset the index to turn it into a normal dataframe:

Doing this resets the index and returns the following:

	Region	Sales
0	East	408.182482
1	North	438.924051
2	South	432.956204
3	West	452.029412

Filtering Python Pivot Tables



Let's filter our data! Source: Nik Piepenbreier

Now that we've created our first few pivot tables, let's explore how to filter the data. Let's create a dataframe that generates the **mean Sale** price by **Region**:

The values in this dataframe are:

Sales	Region
-------	--------

East	408.182482
North	438.924051
South	432.956204
West	452.029412

Now, say we wanted to filter the dataframe to only include **Regions** where the average sale price was over 450, we could write:

Sales	Region
West	452.029412

We can also apply multiple conditions, such as filtering to show only **sales** greater than 450 or less than 430. I know this is a strange example, but it's just illustrative:

Notice that we wrapped each condition in brackets and separated the conditions by a **pipe** (|) symbol. This returns the following:

Sales	Region
East	408.182482
West	452.029412

Creating a Multi-Index Pivot Table

Single index pivot tables are great for generating high-level overviews. However, we can also add additional indices to a pivot table to create further groupings. Say,

we wanted to calculate the sum per Type and per Region, we could write the following:

This will give us the following:

		Sales
Region	Type	
East	Children's Clothing	45849
	Men's Clothing	51685
	Women's Clothing	70229
North	Children's Clothing	37306
	Men's Clothing	39975
	Women's Clothing	61419
South	Children's Clothing	18570
	Men's Clothing	18542
	Women's Clothing	22203
West	Children's Clothing	20182
	Men's Clothing	19077
	Women's Clothing	22217

We could also apply multiple functions to our pivot table. Say we wanted to have the same pivot table that showed us the total sum but also the count of sales, we could write:

		sum	len
		Sales	Sales
Region	Type		
East	Children's Clothing	45849	113
	Men's Clothing	51685	122
	Women's Clothing	70229	176

North	Children's Clothing	37306	85
	Men's Clothing	39975	89
	Women's Clothing	61419	142
South	Children's Clothing	18570	45
	Men's Clothing	18542	39
	Women's Clothing	22203	53
West	Children's Clothing	20182	42
	Men's Clothing	19077	41
	Women's Clothing	22217	53

Adding Columns to a Pandas Pivot Table

Adding columns to a pivot table in Pandas can add another dimension to the tables. Based on the description we provided in our earlier section, the Columns parameter allows us to add a key to aggregate by. For example, if we wanted to see the number of units sold by Type and by Region, we could write:

Let’s visualize what this looks like:

Region	East	North	South	West
Type				
Children's Clothing	2318.0	1763.0	1017.0	78
Men's Clothing	2420.0	0.0	725.0	82
Women's Clothing	3372.0	2596.0	1056.0	10

This allows us to see the data in a different, more easy-to-read format. It also makes plotting the data out a little

easier, as we'll explore in the next section.

Columns might be one of the more confusing parts of the pivot table function, especially with how they relate to values. Columns are optional as we indicated above and provide the keys by which to separate the data. The pivot table aggregates the values in the values parameter.


Plotting Pandas Pivot Tables



Let's plot our data! Source: Nik Piepenbreier

While pivot tables make it easy to summarize data, plotting this data out makes it even easier to identify trends. Let's, for example, create bar plots for the pivot table we generated with columns. We can do this by writing the following:

This generates the following plot:



Graphing our pivot table. Source: Nik Piepenbreier

We can now visualize that the East region has the highest sales across Types, and that the South region has the lowest sales across Types.

If you want to find ways to make Matplotlib charts a little less ugly, check out my other tutorial here:

[Visualizing COVID-19 Data Beautifully in Python \(in 5 Minutes or Less!!\)](#)

[Making Matplotlib a Little Less Painful!](#)

Handling Missing Data in Python Pivot Tables

No data set is perfect! Let's see how we can handle missing data in Python pivot tables. To see which columns have missing data, we can run the info() function to explore the data set:

This returns the following output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        1000 non-null   datetime64[ns]
 1   Region      1000 non-null   object
 2   Type        1000 non-null   object
 3   Units       911 non-null    float64
 4   Sales       1000 non-null   int64
dtypes: datetime64[ns](1), float64(1), int64(1),
memory usage: 39.2+ KB
```

We can see that Units is the only field with missing values. Let's see what happens when we generate this following pivot table:

This returns the following pivot table:

Region	East	North	South	West
Type				
Children's Clothing	35.0	35.0	35.0	35.0
Men's Clothing	35.0	NaN	34.0	35.0
Women's Clothing	35.0	35.0	35.0	35.0

The NaN displayed in the table, of course, doesn't look great. We can choose to fill that value with something else. For example, if we wanted to fill N/A for any missing values, we could write the following:

The output of this is:

Region	East	North	South	West
Type				
Children's Clothing	35.0	35	35.0	35
Men's Clothing	35.0	N/A	34.0	35
Women's Clothing	35.0	35	35.0	35

Much better!

Adding Totals for Rows and Columns to Pandas Pivot Tables

For our last section, let's explore how to add totals to both rows and columns in our Python pivot table. We do this with the margins and margins_name parameters. If we wanted to add this to the pivot table we created above, we would write the following:

This returns the following:

Region	East	North	South	West
Type				
Children's Clothing	35	35	35	35
Men's Clothing	35	N/A	34	35

Women's Clothing	35	35	35	1
Total	35	35	35	1

The margins parameter requires a boolean (True/False) value to either add row/column totals or not. The margins_name parameter allows us to add labels to these values.

Conclusion: Python Pivot Tables — The Ultimate Guide



Thanks so much for reading! Source: Nik Piepenbreier

In this post, we explored how to easily generated a pivot table off of a given dataframe using Python and Pandas.

Pivot tables in Python allow you to easily generate insights into data sets, whether large or small. The multitude of parameters available in the `pivot_table` function allows for a lot of flexibility in how data is analyzed. In this post, we explored how to generate a pivot table, how to filter pivot tables in Python, how to add multiple indices and columns to pivot tables, how to plot pivot tables, how to deal with missing values, and how to add row and column totals.

Originally published at <https://datagy.io> on April 12, 2020.